


The Interpretability Challenge in Neural Networks

Model Interpretability Spectrum



Why Interpretability Matters:

1. Trust: Users need to understand predictions
2. Debugging: Find and fix errors
3. Compliance: Regulations (GDPR, finance)
4. Fairness: Detect bias in decisions
5. Science: Gain insights from models

Making NNs More Interpretable

Feature Importance

Which inputs matter most?

SHAP Values

Attribution per feature

LIME

Local linear approximations

Attention Visualization

What the model "looks at"

Gradient-based Methods

Sensitivity analysis

Trade-off: More interpretable models often have lower accuracy
Choose based on application requirements!